

Understanding Experimental Data

Eric Grimson

MIT Department Of Electrical Engineering and
Computer Science

Announcements

- Reading: Chapter 18
- No lecture on Wednesday

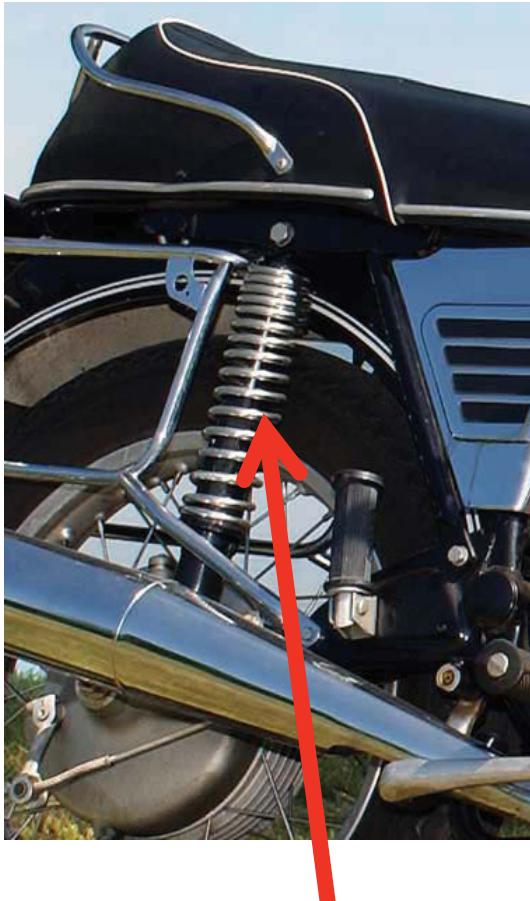
Statistics Meets Experimental Science

- Conduct an experiment to gather data
 - Physical (e.g., in a biology lab)
 - Social (e.g., questionnaires)
- Use theory to generate some questions about data
 - Physical (e.g., gravitational fields)
 - Social (e.g., people give inconsistent answers)
- Design a computation to help answer questions about data

Net Gain on a missed jump shot

$$= P(\text{off reb}) \times E[\text{pts for}] - P(\text{def reb}) \times E[\text{pts against}]$$
- Consider, for example, a spring

This Kind of Spring



$$k \approx 35,000 \text{ N/m}$$

$$k \approx 1 \text{ N/m} \rightarrow$$



Linear spring: amount of force needed to stretch or compress spring is linear in the distance the spring is stretched or compressed

Each spring has a spring constant, k , that determines how much force is needed

Newton = force to accelerate 1 kg mass 1 meter per second per second

Hooke's Law

- $F = -kd$
- How much does a rider have to weigh to compress spring 1cm?

$$F=0.01m*35,000N/m$$

$$F=350N$$

$$mass*9.8m/s^2 = 350N$$

$$mass=350N/9.81m/s^2$$

$$mass=350k/9.81 \leftarrow \text{This } k \text{ refers to kilograms, not the spring constant!}$$

$$mass \approx 35.68k$$

$$F=mass*acc$$

$$F=mass*9.8m/s^2$$



Images of suspension spring © source unknown. All rights reserved.
This content is excluded from our Creative Commons license. For more information, see <https://ocw.mit.edu/help/faq-fair-use/>.

Finding k

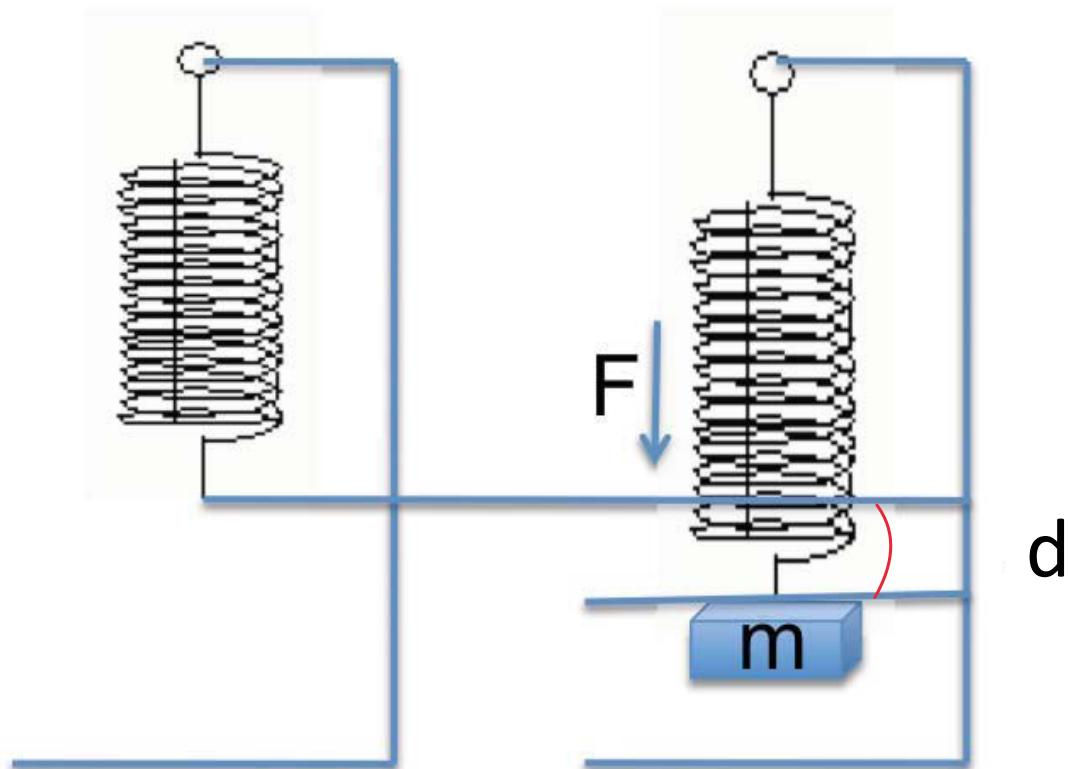
- $F = -kd$
- $k = -F/d$
- $k = 9.81 * m/d$

$$P) F = -kd$$

$$= F \cdot \frac{1}{d} = -k \cancel{d} \cdot \frac{1}{\cancel{d}}$$

$$= \frac{F}{d} = -k$$

$$\Rightarrow k = -\frac{F}{d}$$



Some Data

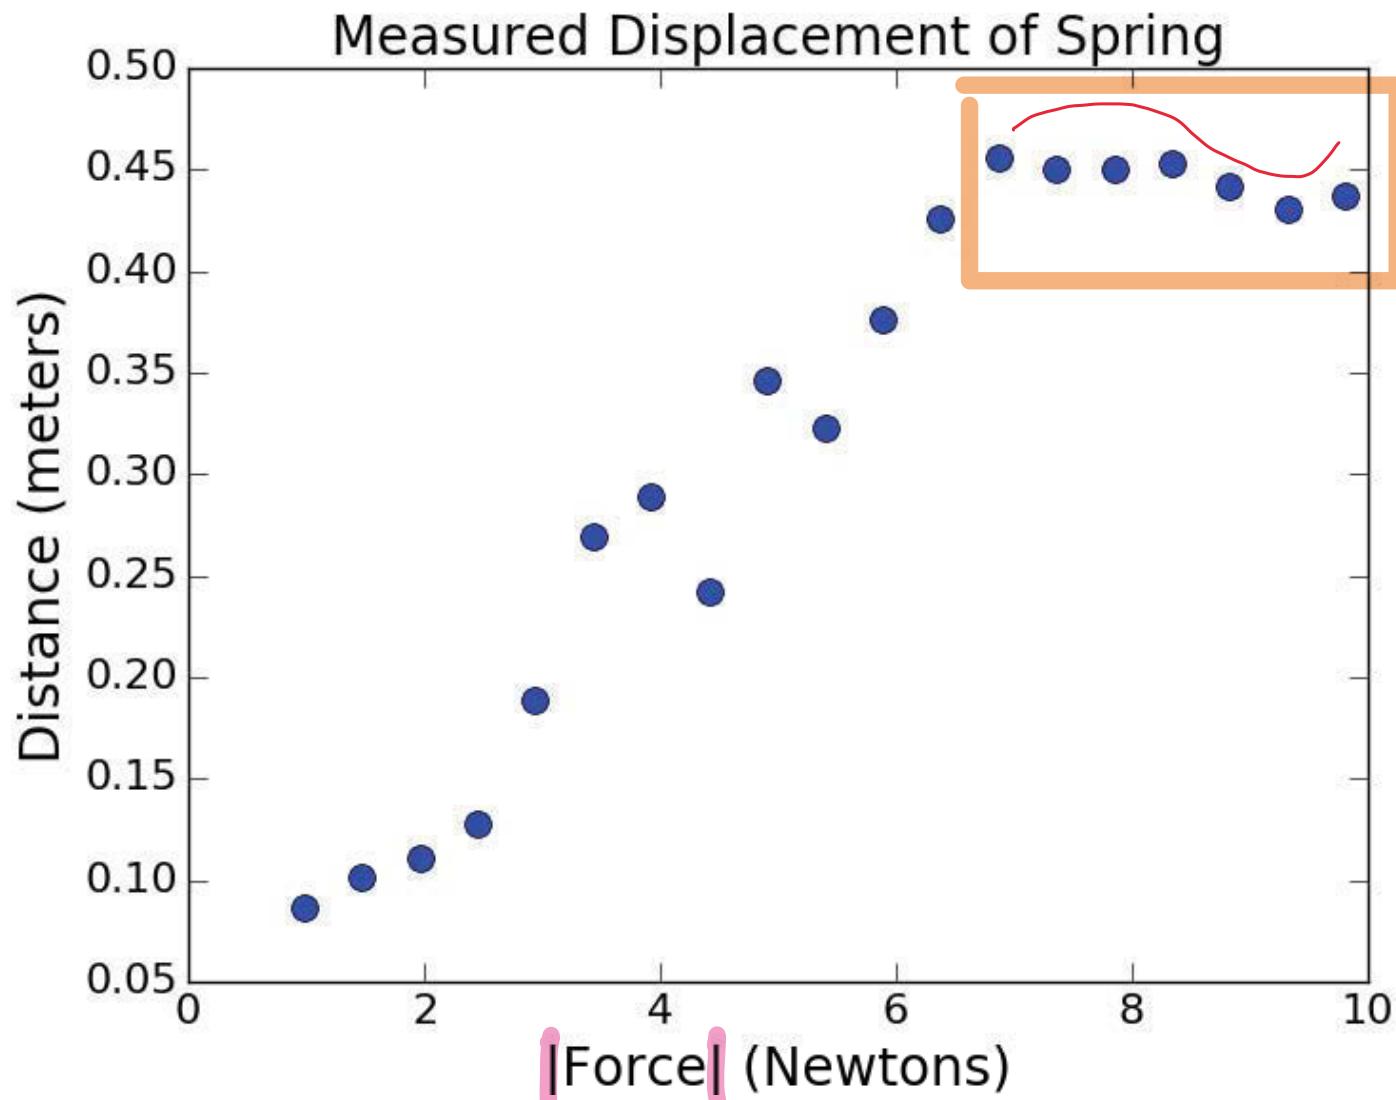
| Distance (m) | Mass (kg) |
|--------------|-----------|
| 0.0865 | 0.1 |
| 0.1015 | 0.15 |
| 0.1106 | 0.2 |
| 0.1279 | 0.25 |
| 0.1892 | 0.3 |
| 0.2695 | 0.35 |
| 0.2888 | 0.4 |
| 0.2425 | 0.45 |
| 0.3465 | 0.5 |
| 0.3225 | 0.55 |
| 0.3764 | 0.6 |
| 0.4263 | 0.65 |
| 0.4562 | 0.7 |

Taking a Look at the Data

```
def plotData(fileName):
    xVals, yVals = getData(fileName)
    xVals = pylab.array(xVals)
    yVals = pylab.array(yVals)
    xVals = xVals*9.81 #acc. due to gravity
    pylab.plot(xVals, yVals, 'bo',
               label = 'Measured displacements')
labelPlot()
```

From
Notes.

Taking a Look at the Data

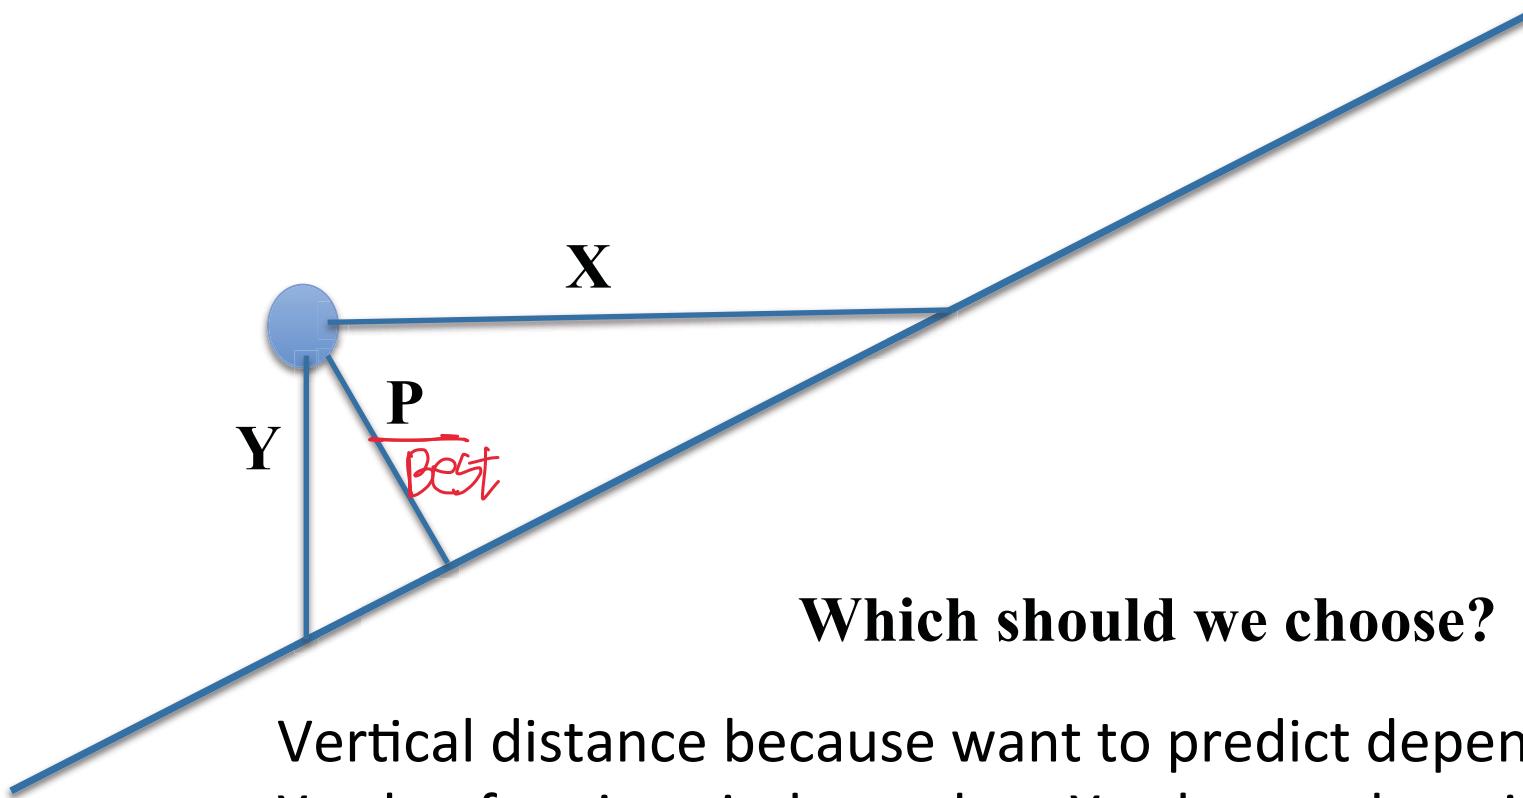


Always $F > 0$

Fitting Curves to Data

- When we fit a ~~curve to a set of data~~, we are ~~finding a~~ fit that relates an independent variable (the mass) to an estimated value of a dependent variable (the distance)
- To decide how well a curve fits the data, we need a way to measure the goodness of the fit – called the **objective function**
- Once we define the objective function, we want to find the curve that minimizes it
- In this case, we want to find a line such that some function of the sum of the distances from the line to the measured points is minimized

Measuring Distance



Vertical distance because want to predict dependent Y value for given independent X value, and vertical distance measures error in that prediction

Least Squares Objective Function

$$\sum_{\substack{i=0 \\ \text{cur index.}}}^{\text{len(observed)}-1} (\text{observed}[i] - \text{predicted}[i])^2$$

Always > 0

- Look familiar?
 - This is variance times number of observations
 - So minimizing this will also minimize the variance

Solving for Least Squares

$$\sum_{i=0}^{\text{len}(\text{observed})-1} (\text{observed}[i] - \text{predicted}[i])^2$$

- To minimize this objective function, want to find a curve for the predicted observations that leads to minimum value
- Use **linear regression** to find a polynomial representation for the predicted model

Polynomials with One Variable (x)

- 0 or sum of finite number of non-zero terms
- Each term of the form cx^p Ex) $\underline{4}x^{\underline{2}} - p$
 - c, the coefficient, a real number c
 - p, the degree of the term, a non-negative integer
- The degree of the polynomial is the largest degree of any term

■ Examples

◦ Line: $ax + b$ $y = f(x)$

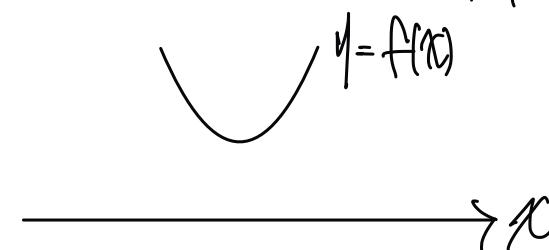
◦ Parabola: $ax^2 + bx + c$

$$D = b^2 - 4ac$$

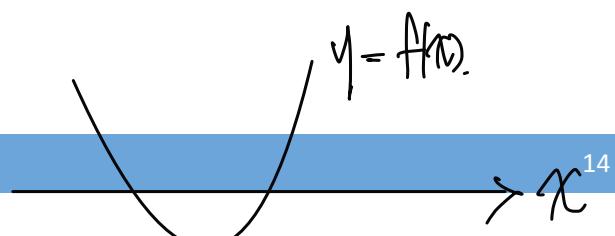
$$D = 0$$



$$D > 0$$



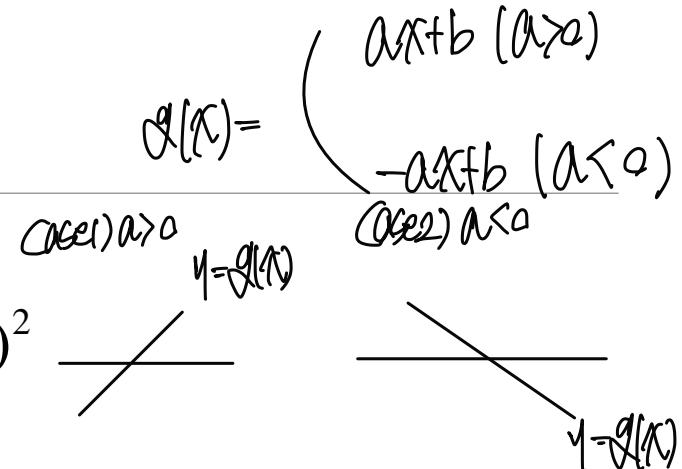
$$D < 0$$



Solving for Least Squares

$\text{len}(\text{observed}) - 1$

$$\sum_{i=0}^{\text{len}(\text{observed})-1} (\text{observed}[i] - \text{predicted}[i])^2$$



- Simple example:

- Use a degree-one polynomial, $y = ax + b$, as model of our data (we want best fitting line)

- Find values of a and b such that when we use the polynomial to compute y values for all of the x values in our experiment, the squared difference of these predicted values and the corresponding observed values is minimized

- A **linear regression** problem

- Many algorithms for doing this, including one similar to Newton's method (which you saw in 6.0001)

polyFit

- Good news is that pylab provides built in functions to find these polynomial fits
- pylab.polyfit(observedX, observedY, n)
- Finds coefficients of a polynomial of degree n, that provides a best least squares fit for the observed data
 - $n = 1$ – best line $y = ax + b$
 - $n = 2$ – best parabola $y = ax^2 + bx + c$

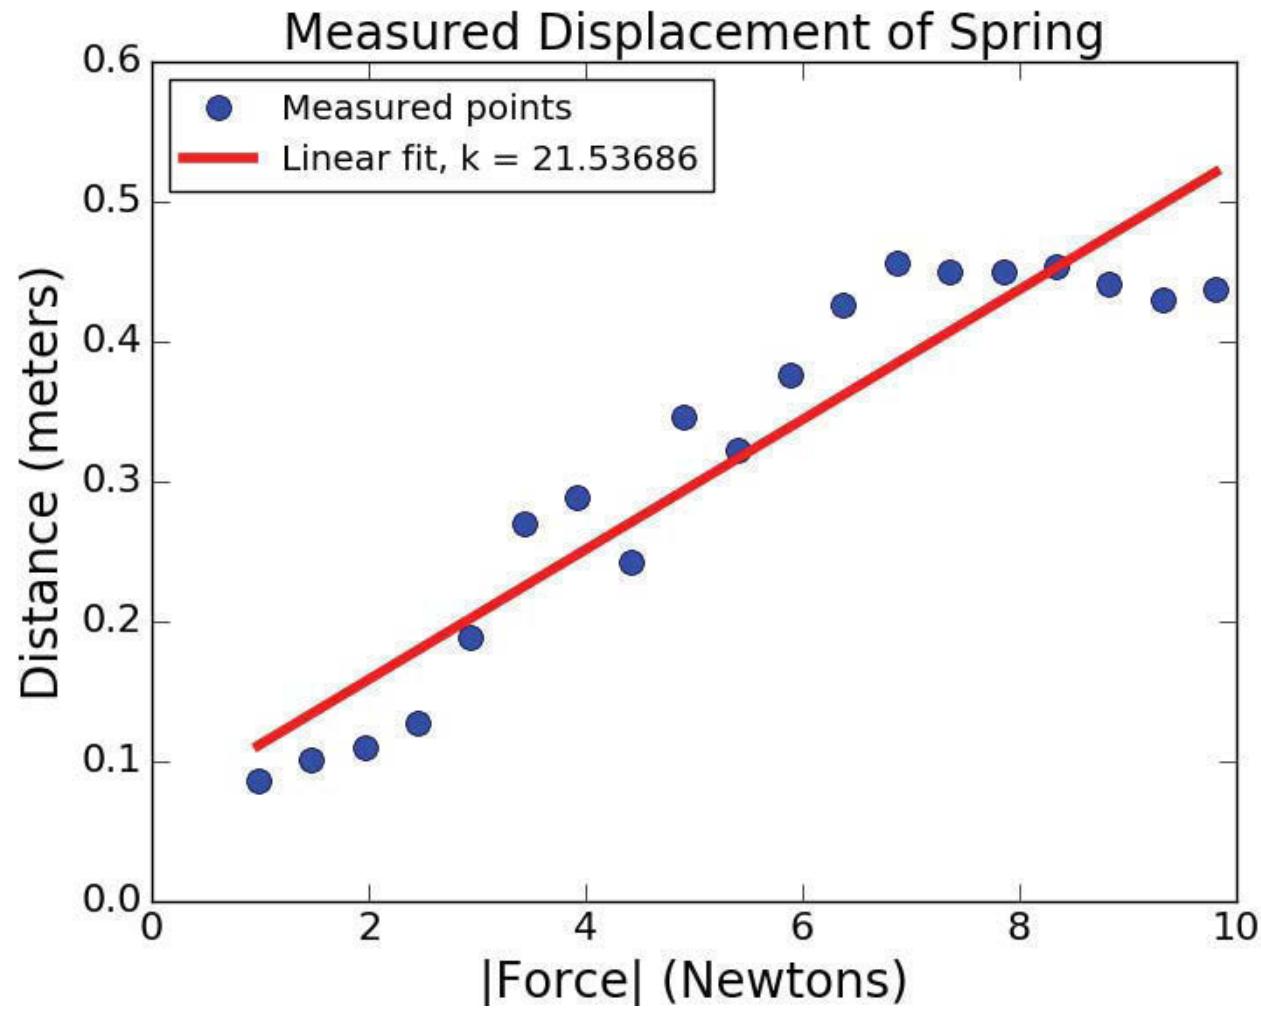
Using polyfit

```
def fitData(fileName):
    xVals, yVals = getData(fileName)
    xVals = pylab.array(xVals)
    yVals = pylab.array(yVals)
    xVals = xVals*9.81 #get force
    pylab.plot(xVals, yVals, 'bo',
               label = 'Measured points')
    labelPlot()
    a,b = pylab.polyfit(xVals, yVals, 1)
    estYVals = a*pylab.array(xVals) + b
    print('a = ', a, 'b = ', b)
    pylab.plot(xVals, estYVals, 'r',
               label = 'Linear fit, k = '
               + str(round(1/a, 5)))
    pylab.legend(loc = 'best')
```

plotData

Note that conversion to array is redundant here

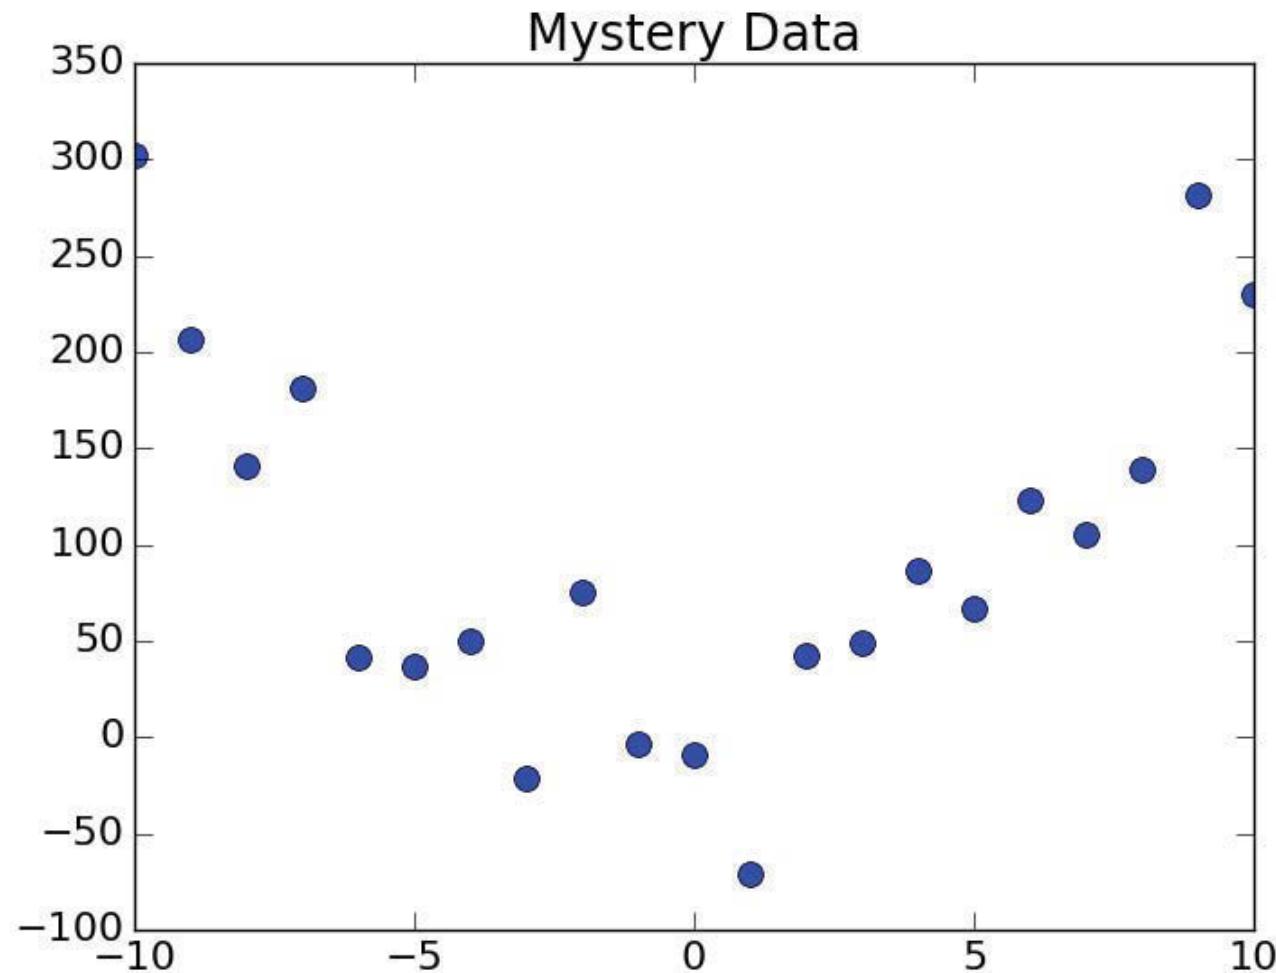
Visualizing the Fit



Version Using polyval

```
def fitData1(fileName):
    xVals, yVals = getData(fileName)
    xVals = pylab.array(xVals)
    yVals = pylab.array(yVals)
    xVals = xVals*9.81 #get force
    pylab.plot(xVals, yVals, 'bo',
                label = 'Measured points')
    labelPlot()
    model = pylab.polyfit(xVals, yVals, 1)
    estYVals = pylab.polyval(model, xVals)
    pylab.plot(xVals, estYVals, 'r',
                label = 'Linear fit, k = '
                + str(round(1/model[0], 5)))
    pylab.legend(loc = 'best')
```

Another Experiment



Fit a Line

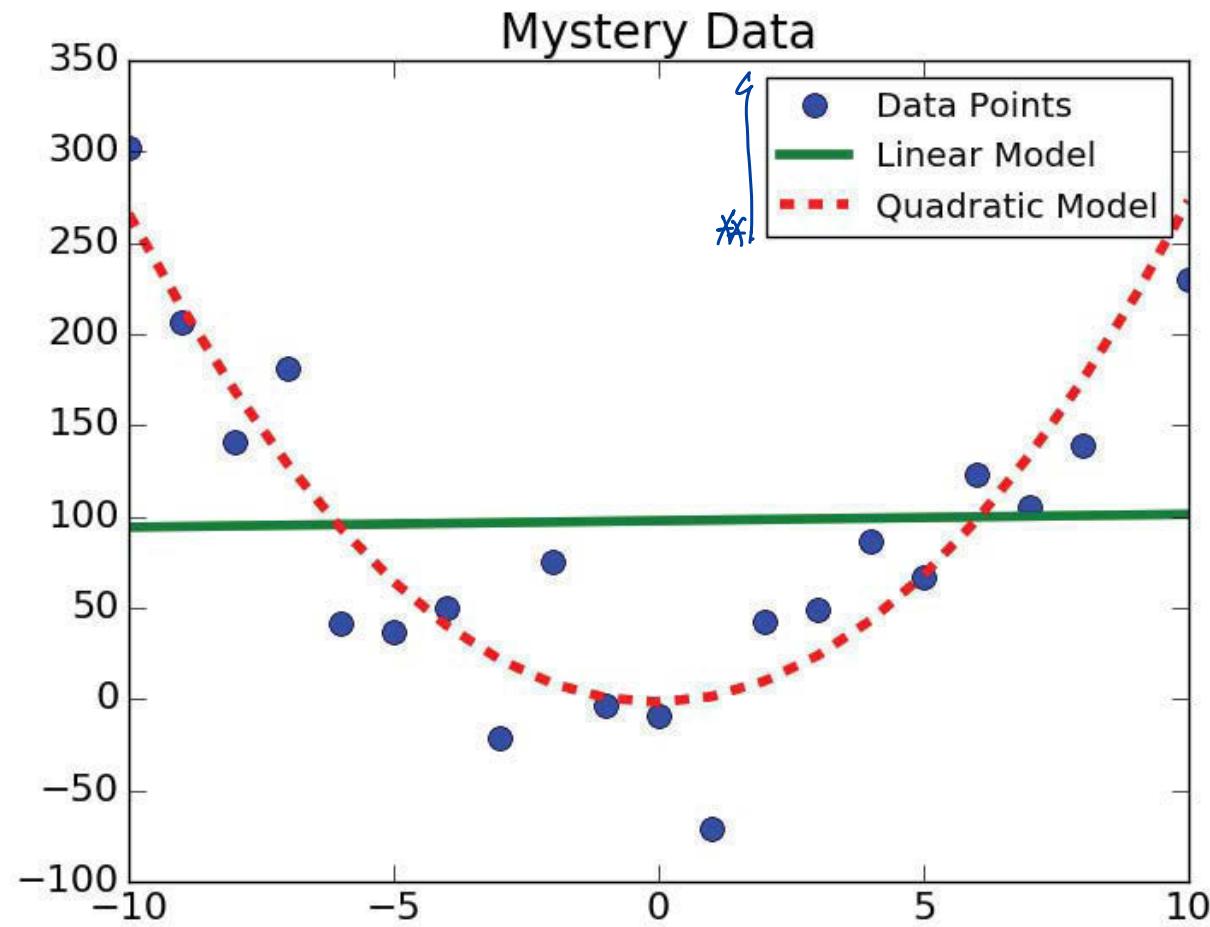


Let's Try a Higher-degree Model

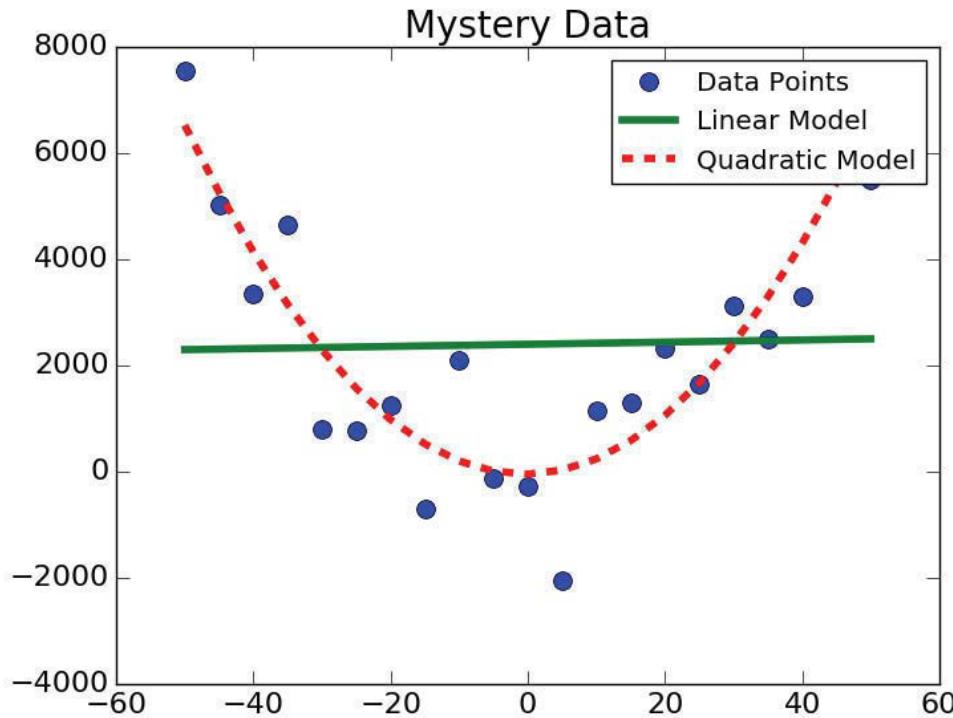
```
model2 = pylab.polyfit(xVals, yVals, 2)
pylab.plot(xVals, pylab.polyval(model2, xVals),
           'r--', label = 'Quadratic Model')
```

Note that this is still an example of linear regression,
even though we are not fitting a line to the data (in this
case we are finding the best parabola)

Quadratic Appears to be a Better Fit



How Good Are These Fits?



- Relative to each other
- In an absolute sense

Relative to Each Other

- Fit is a function from the independent variable to the dependent variable
- Given an independent value, provides an estimate of the dependent value
- Which fit provides better estimates?
- Since we found fit by minimizing mean square error, could just evaluate goodness of fit by looking at that error

Comparing Mean Squared Error

```
def aveMeanSquareError(data, predicted):
    error = 0.0
    for i in range(len(data)):
        error += (data[i] - predicted[i])**2
    return error/len(data)

estYVals = pylab.polyval(model1, xVals)
print('Ave. mean square error for linear model =',
      aveMeanSquareError(yVals, estYVals))
estYVals = pylab.polyval(model2, xVals)
print('Ave. mean square error for quadratic model =',
      aveMeanSquareError(yVals, estYVals))
```

Ave. mean square error for linear model = 9372.73078965

Ave. mean square error for quadratic model = 1524.02044718

In an Absolute Sense

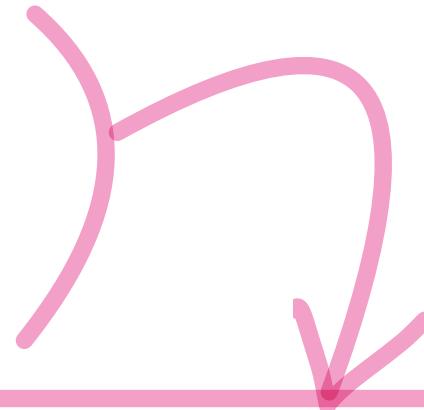
- Mean square error useful for comparing two different models for the same data
- Useful for getting a sense of absolute goodness of fit?
 - Is 1524 good?
- Hard to know, since there is no upper bound and not scale independent
- Instead we use coefficient of determination, R^2 ,

$$R^2 = 1 - \frac{\sum_i (y_i - p_i)^2}{\sum_i (y_i - \mu)^2}$$

y_i are measured values \leftarrow Error in estimates
 p_i are predicted values \leftarrow Variability in
 μ is mean of measured values measured data

If You Prefer Code

$$R^2 = 1 - \frac{\sum_i (y_i - p_i)^2}{\sum_i (y_i - \mu)^2}$$



```
def rSquared(observed, predicted):  
    error = ((predicted - observed)**2).sum()  
    meanError = error/len(observed)  
    return 1 - (meanError/numpy.var(observed))
```

I am playing a clever trick here:

- Numerator is sum of squared errors
- Dividing by number of samples gives average sum-squared-error
- Denominator is variance times number of samples
- So mean SSE/variance is same as R^2 ratio

R²

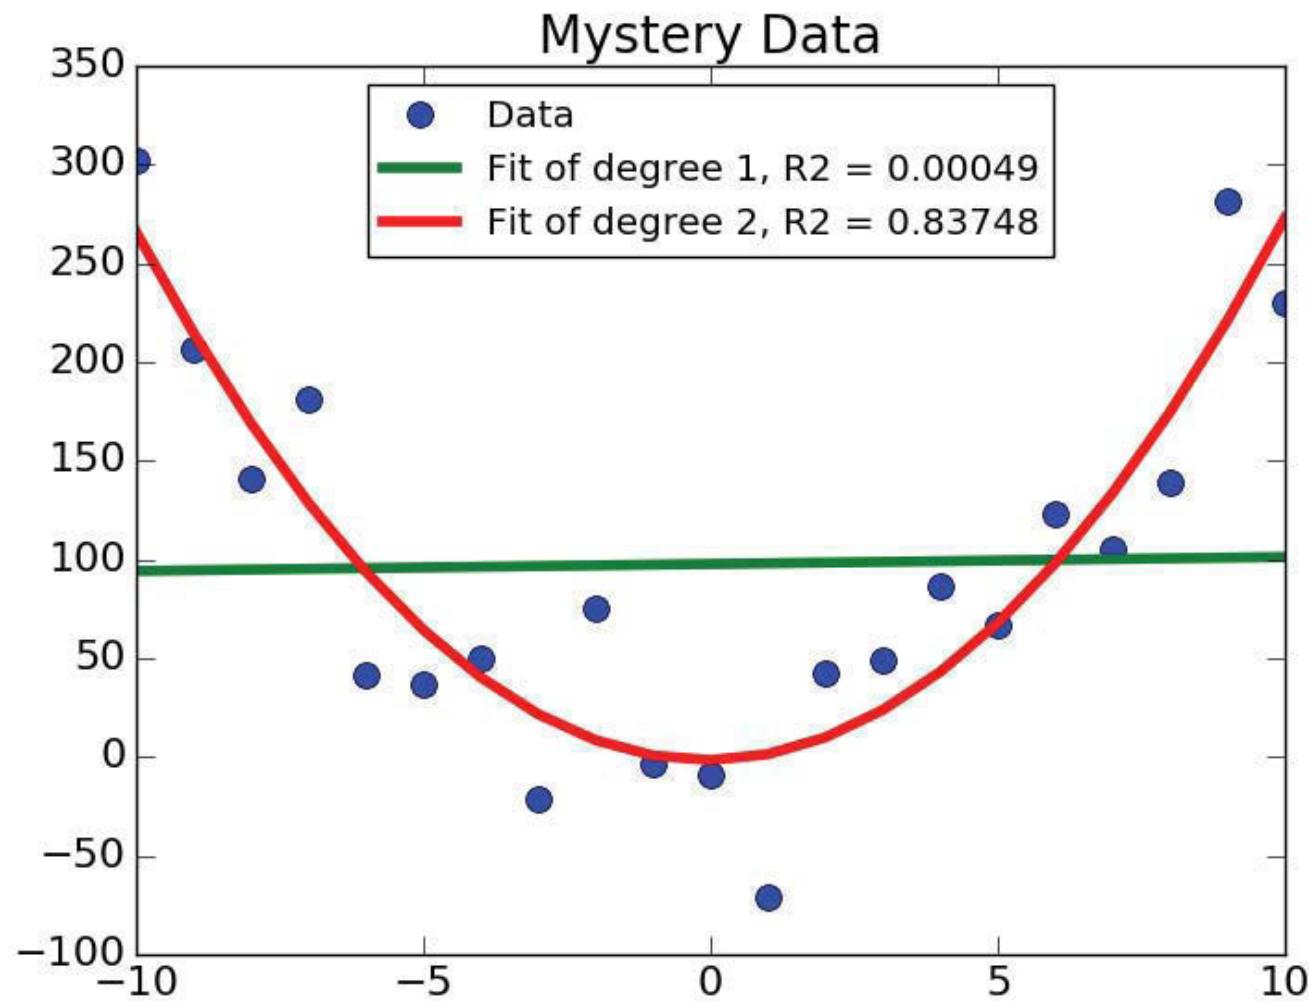
- By comparing the estimation errors (the numerator) with the variability of the original values (the denominator), R² is intended to capture the proportion of variability in a data set that is accounted for by the statistical model provided by the fit
- Always between 0 and 1 when fit generated by a linear regression and tested on training data
 - If R² = 1, the model explains all of the variability in the data.
 - If R² = 0, there is no relationship between the values predicted by the model and the actual data.
 - If R² = 0.5, the model explains half the variability in the data.

Testing Goodness of Fits

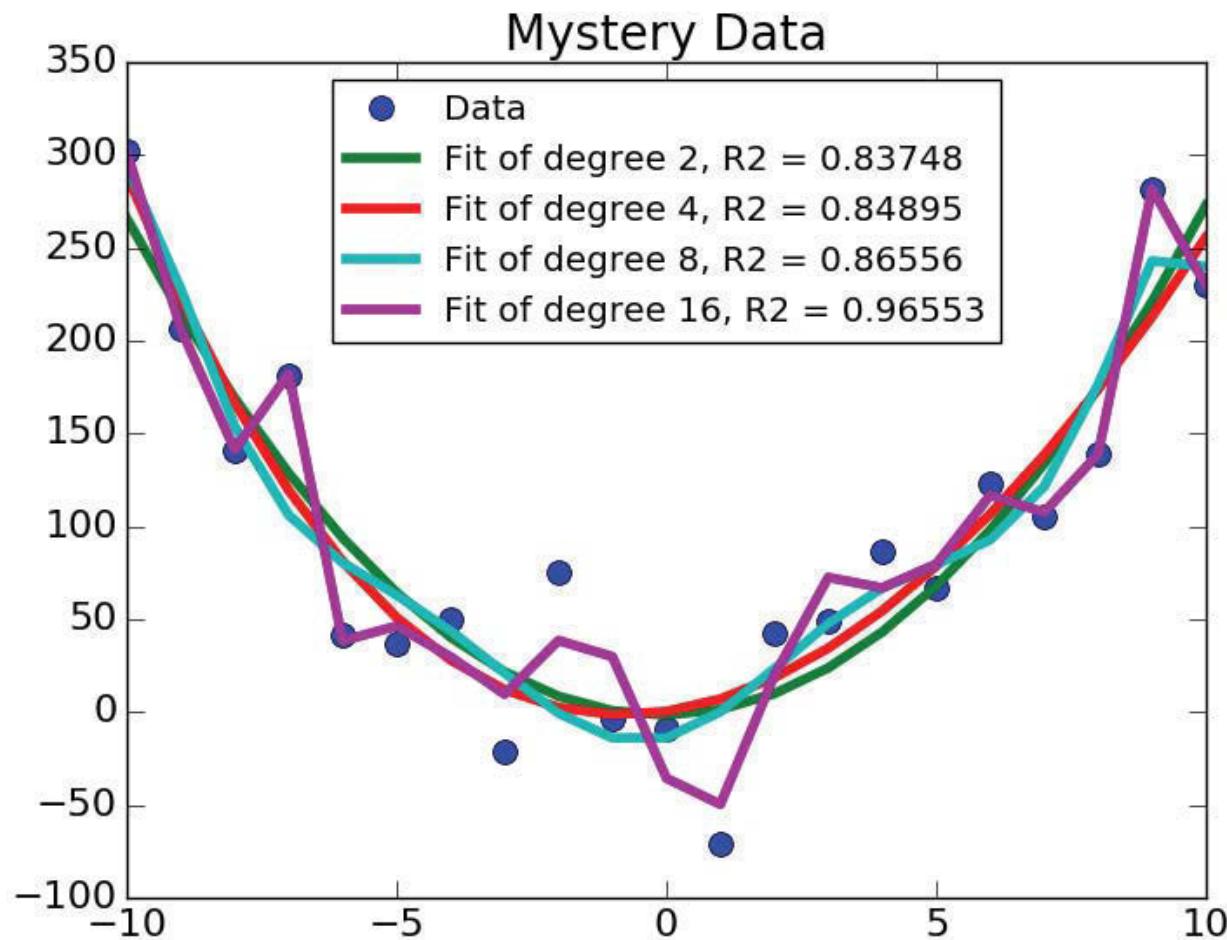
```
def genFits(xVals, yVals, degrees):
    models = []
    for d in degrees:
        model = pylab.polyfit(xVals, yVals, d)
        models.append(model)
    return models

def testFits(models, degrees, xVals, yVals, title):
    pylab.plot(xVals, yVals, 'o', label = 'Data')
    for i in range(len(models)):
        estYVals = pylab.polyval(models[i], xVals)
        error = rSquared(yVals, estYVals)
        pylab.plot(xVals, estYVals,
                   label = 'Fit of degree ' \
                   + str(degrees[i]) \
                   + ', R2 = ' + str(round(error, 5)))
    pylab.legend(loc = 'best')
    pylab.title(title)
```

How Well Fits Explain Variance



Can We Get a Tighter Fit?



MIT OpenCourseWare

<https://ocw.mit.edu>

6.0002 Introduction to Computational Thinking and Data Science

Fall 2016

For information about citing these materials or our Terms of Use, visit: <https://ocw.mit.edu/terms>.